

Kanishk Hari

(949) 933-8652 | kanishkharicollege@gmail.com | Irvine, CA, USA | [linkedin.com/in/kanishk-hari/](https://www.linkedin.com/in/kanishk-hari/) | github.com/KanishkHari

EDUCATION

University of California – San Diego

Bachelor of Science in Data Science

Graduating June 2026

GPA: 3.5/4.0

Relevant Coursework: Introduction to Artificial Intelligence/Machine Learning, Object-Oriented Programming, Probability & Statistics, Data Structures & Algorithms, Web Mining and Recommender Systems, Applications of Data Science

WORK EXPERIENCE

Siloq

Associate Data Scientist

June 2025 – Present

Irvine, CA

- Enhanced learning for 3,500+ students by developing an AI-powered study assistant using RAG, Go, and the GPT-3.5 Turbo API, improving content retrieval speed by 200% for faster question resolution.
- Streamlined release cycles by deploying data-driven microservices on AWS Elastic Kubernetes Service (EKS) with automated CI/CD pipelines using GitHub Actions, prompting a 37% reduction in production errors.
- Analyzed user engagement metrics following the launch of a custom QR tool, correlating accessibility improvements to a 54% increase in traffic and \$1,500+ in revenue.

National Science Foundation (UC Irvine)

Data Science Research Fellow

June 2024 – July 2024

Irvine, CA

- Engineered end-to-end data pipeline for 69K+ earthquake and bridge records using Python, R, and Scikit-learn, enabling large-scale analysis of seismic collapse risk.
- Optimized LASSO and GAM models to predict column drift ratios with 93% accuracy, improving adjusted R^2 by 15% through feature selection and hyperparameter tuning.
- Delivered statistically validated insights to 100+ researchers, informing design recommendations that enhance bridge resilience under seismic conditions.

California State University, Fullerton

Undergraduate Research Assistant

May 2023 – July 2023

Fullerton, CA

- Designed and trained symbolic regression models in Python to automatically derive mathematical expressions from datasets, achieving 90% accuracy on benchmark tasks.
- Enhanced model efficiency through vectorization and hyperparameter tuning, increasing execution speed by 25% and reducing development overhead by 40%.
- Validated algorithm performance on synthetic and empirical datasets, ensuring generalization through k-fold cross-validation and performance benchmarking.

PROJECTS

Content-Based Movie Recommendation Engine | Python, Pandas, Scikit-learn, NLP

- Built a content-based recommendation engine using NLP for 8,800+ Netflix titles utilizing TF-IDF vectorization and Cosine Similarity to extract semantic importance from unstructured plot descriptions, creating a high-dimensional feature space for similarity analysis.
- Developed a query-efficient retrieval algorithm that ranks and returns top-N content matches in real-time, translating high-dimensional vector space distances into actionable viewer suggestions, increasing user engagement by 18% during testing with simulated viewing sessions.

Modeling Hidden Pollution Regimes with HMMs | Python, NumPy, Scikit-learn

- Developed a two-state Hidden Markov Model (HMM) to infer unobserved weather regimes from air quality data, achieving a 0.68 F-score by implementing Viterbi decoding and Forward-Filtering Backward-Sampling in Python, enabling accurate prediction of high-risk air pollution days for city planning simulations.
- Engineered a synthetic labeling pipeline for rainfall detection by mapping EPA Air Quality Index (AQI) breakpoints to pollutant concentrations (CO , NO_2 , O_3), enabling model evaluation on a previously unlabeled UCI dataset and informing environmental policy decisions.

TECHNICAL SKILLS

Languages: Python, R, SQL, C/C++, Go, Java, JavaScript

Libraries & Tools: Pandas, NumPy, Scikit-learn, PyTorch, TensorFlow, Matplotlib, Seaborn, Tableau, Git, Power BI, Spark, Dask

Databases & Cloud: AWS, Postgres, MySQL, NoSQL, REST API, MongoDB

Concepts: Statistical Inference, Machine Learning, Predictive Modeling, Feature Engineering, Data Cleaning, Visualization, A/B Testing, Data Pipelines, ETL Pipelines, Data Lake, Data Warehouse