Candidate Name

Phone: (XXX) XXX-XXXX · Email: firstname.lastname@email.com · LinkedIn: linkedin.com/in/username · GitHub: github.com/username

Education

Large Public Research University — City, State
Master of Science - Computer Science (GPA: 3.5/4.0) — **Aug 2024 – May 2026**
- Coursework: Network Systems, Big Data Architecture, Software Engineering, Natural Language Processing

Private Engineering University — City, Country
Bachelor of Engineering - Information Science & Engineering (GPA: 3.8/4.0) — **Aug 2019 – May 2023**
- Coursework: Machine Learning, Statistics and Probability, Cloud Computing, Distributed Computing

Experience

Graduate Research Assistant — City, State
Research Lab, Large Public Research University — **Jun 2025 – Present**
- Developed an ML pipeline for **markerless 3D pose estimation** in **TensorFlow**, automating preprocessing and kinematic feature extraction to process millions of frames and generate standardized motion features, achieving a **30%** faster throughput.
- Engineered a hierarchical **Video Swin Transformer** architecture employing multi-scale temporal attention for real-time classification of complex behavioral motifs, achieving $>$ **90% frame-wise F1-score** and reducing manual labeling effort by over **75%**.

Software Engineer — City, State
University Research Center — **Feb 2025 – Present**
- Developed **flight-critical software** for an NSF-funded small-satellite mission, achieving launch readiness of three satellites for low Earth orbit.
- Cut false alarms by **15%** and improved resiliency by enhancing deployment and safety logic; built distributed **Python** pipelines using **Redis** for sub-ms caching and **AWS DynamoDB** for high-concurrency storage handling **50K+ ops/sec**.

Software Engineer Intern — Remote
B2B SaaS Startup — **May 2025 – Aug 2025**
- Productionized **LLM**-powered calendar intelligence by building **Go** microservices with **REST** and gRPC APIs for scheduling and categorization, shipping to **production** with **CI/CD** and real-time monitoring, establishing metrics, and leading technical and architectural decisions.
- Improved uptime to **99.99%** and reduced API latency by **40%** by eliminating redundant database queries, introducing in-memory caching, optimizing Go hot paths with concurrency, migrating endpoints to **gRPC**, and deploying load balancing with autoscaling for traffic spikes.

Software Engineer — City, Country
Fintech Company — **Jul 2023 – Jun 2024**
- Directed **Java** microservices and **AWS Lambda + RDS** ingestion services for data export; doubled throughput and cut server costs by **25%**.
- Spearheaded a team of 3 to develop and modernize policy systems, shipped **9** multi-threaded APIs using an MVP approach, and built **Flink** pipelines on **Kafka** and **AWS SQS** to handle **1M** events/day with low-latency, reliable processing.

Project Intern — City, Country
Fintech Company — **Feb 2023 – Jul 2023**
- Built **Spring Boot** microservices and automated **ETL** for ingesting insurance claims and policy data to produce standardized reports, cutting runtime by **35%** through parallelism. Delivered **React** and **Redux** dashboards with job status and searchable history, reducing support requests by **50%**.
- Optimized **Jenkins CI/CD** pipelines and standardized **Gradle** builds, cutting bug-fix turnaround by **20%**. Improved system security by reducing access incidents by **18%** through scoped tokens and MFA, contributing to a **9%** increase in overall user satisfaction based on survey.

Projects

- Research Paper QA & Synthesis Platform | *FastAPI, LangChain, Pinecone, Python, LLM APIs, Docker*
  - Developed a **Retrieval-Augmented Generation (RAG)** pipeline using **FastAPI** and **LangChain** to turn long research papers into searchable data stored in **Pinecone**, enabling question answering across thousands of documents.
  - Developed an automated synthesis engine leveraging **LLM-based reasoning** to perform cross-document analysis, generating tailored research summaries and identifying knowledge gaps across thousands of papers to reduce literature review time by **45%**.
- Observability Tooling – Open Source Contribution | *Go, Shell*
  - Standardized codebase health by architecting a unified linting and formatting framework, enforcing strict consistency via **pre-commit hooks** and **automated validation checks** across a high-traffic distributed repository.
  - Engineered **auto-fix CI pipelines** that resolved **10% of legacy linting issues** and increased code review turnaround by **15%**, reducing manual developer workload and shifting error detection earlier in development.

Technical Skills

- **Languages:** Python, Java, Go, C++, JavaScript, TypeScript, SQL, Shell
- **Frameworks:** Spring Boot, React, Node.js, Django, Flask, FastAPI, Flink
- **Databases:** AWS DynamoDB, PostgreSQL, MongoDB, Pinecone, Redis
- **Tools & Platforms:** AWS (EC2, S3, RDS, Lambda, CloudFront), Docker, Kubernetes, Terraform, Git, Jenkins, Kafka, Linux
- **ML & AI:** PyTorch, JAX, TensorFlow, OpenCV, Transformers/LLMs, RAG, LangChain
- **Concepts:** Machine Learning, Data Structures & Algorithms, Object-Oriented Design, Distributed Systems, System Design